

Exercise 2: Reproduce experimental results from a paper

Deadline for submission: Sun, January 24th, 2021, 23:59

For this assignment, you will work in **groups of three** and **reproduce** the experimental setup, experiments, and results as explained in a scientific paper. The task of interest is – based on careful experimental design – to confirm the numbers and findings reported in the paper or, alternatively, uncover inconsistencies and therefore challenge the conclusions drawn.

Specifically, by reproducing the experiments, it can be checked whether

- the information given in the paper is sufficient to reproduce the results reported
- statistically significant differences can be made out between the different settings, in particular when this is not reported in the paper (no significance tests, confidence intervals, p values, etc. or not even variance)
- the values reported in the paper could stem from the distribution sampled by the reproduced experiments

Your efforts need to be documented in a report that should be accompanied by the code and workflow you used to reproduce the experiments. To do so, first, identify the experimental setup and strategy taken in the paper and outline the steps necessary for you to reproduce results. Try to stay to the described steps as close as possible, i.e., use the same implementations and settings wherever possible. If the exact implementation is unknown or unavailable, find a reasonable substitute. You might have to consult additional sources to get the full picture of the used data or methods.

In each step, document which information is given in the paper and which measures you took to implement this step, i.e., which implementation and which parameters you used. Report the numbers obtained in each intermediate step. Identify deviations from the numbers reported in the original paper, their origin, and estimate whether they will have a significant impact on subsequent steps.

For the results, perform adequate tests to test for statistical significance. Justify your choice. Can you confirm the findings of the paper? Did you identify a flaw in their setup? Could you correct the flaw? How did it impact the results and findings of the experiments?

You need to produce the **results of one scientific paper out of three candidates you need to identify and assess**.

In the first step, find three candidate papers across the last 5 years of proceedings of these three conference series:

- SIGIR: International ACM SIGIR Conference on Research and Development in Information Retrieval
- CIKM: International Conference on Information and Knowledge Management
- ECIR: European Conference on Information Retrieval

For SIGIR and CIKM search the proceedings on <https://dl.acm.org> and use filters to restrict your search, for ECIR, proceedings are published by Springer, another way to access them is via dblp (<https://dblp.org/db/conf/ecir/index.html>, “Advances in Information Retrieval” entries).

Examine 3 papers from these conferences for their likelihood regarding reproducibility. Enter into the provided TUWEL form the DOI, author, title and conference for the papers examined (each group should enter different papers, hence register your candidates first by entering them into the list). After examination, add your assessment according to the criteria whether code and dataset used are available (providing dataset URI, platform information, etc. if applicable). Highlight which of your candidates you chose for your reproducibility project.

If the chosen paper is very easily reproducible, i.e. all code is available and runs out of the box, producing exactly the results as described in the paper, focus even more on performing a detailed examination of the experimental setup and put the results under scrutiny: were experiments run multiple times, e.g. using different seeds for data splits; was variance reported; were significance tests performed (applicable ones); do conclusions hold if you rerun experiments using different seeds?

For writing the report, follow the ACM formatting guidelines, using the templates provided at <https://www.acm.org/publications/proceedings-template>. (Proceedings Style File: LaTeX2e - Strict Adherence to SIGS style; LaTeX recommended, but Word/OpenOffice is also ok). **Report page limit: Maximum 6 pages!** Focus on the key aspects!

Upload your report to the TU Wien ExDDS WS 2020/21 community on Zenodo (<https://zenodo.org/deposit/new?c=tuw-exdds-ws20>, choose Report as submission type). Please, make it open or embargoed (embargo end = deadline for this exercise). As a result, you and the other students will have a chance to compare and discuss your findings. When writing your report, keep in mind that it will be publicly visible and open to the research community. For submission in TUWEL, apart from your presentation slides (see below), you will only provide the DOI pointing to your report and experimentation scripts/code/workflows.

In class on Jan 7 or 14, 2021, each group will have to give a short **presentation on the status of the work** (Strict time limit: 5 minutes!). To this end, prepare a deck of about 5 slides (including first slide containing group number, members, and chosen option), presenting your strategies, encountered difficulties and key findings, as well as your conclusion about the experimental design, intermediate results, conclusions and remaining work. Register for a presentation time slot on TUWEL until Jan 6, 2021.

The final report should detail the same aspects: strategies, difficulties, key findings, as conclusion about the experimental design, results, and conclusions.

Grading scheme (*updated*)!:

- Investigation of 3 papers, entries to Registration DB: max. 15%
- Reproduction of results (code and report): max. 40%
- Description, interpretation, statistical testing: max. 25%
- Quality/clarity of report and presentation: max. 15%
- Formal and complete submission to Zenodo: max. 5%